AI-assisted MQ and Relevant Experience Pilot Canvas

Section 2 — Target Audience

- **Name/roles:** Compensation staff, especially Compensation Analysts and the Compensation Manager.
- **Current AI fluency (1–5):** 1–2.
- **What they need:** A practical, structured workflow that helps with complex MQ and relevant experience review while fitting existing resume review work.
- **Likely concern:** That AI will be unreliable, create extra work, or feel risky in a high-stakes HR setting.

Section 1 — Problem Statement

MQ and relevant experience reviews can vary by analyst, especially for complex resumes, with inconsistent documentation and limited bias checks. This creates uneven outcomes and makes decisions harder to explain. A small AI-assisted pilot could improve consistency, documentation, and fairness while keeping final decisions human.

Section 3 — AI Tools and Approach

- **Tool 1:** U of A GenAI with Claude Sonnet 4.6 through Open WebUI.
    - Modality: Text-based AI.
    - Rationale: Institutionally governed environment for structured resume-to-job-description review.
    - Data sensitivity: Resume data high; job descriptions low; prompt medium.
    - Institutional approval: Yes for the platform environment.
- **Tool 2:** Python/Streamlit application calling Claude Sonnet 4.6 through an API.
    - Modality: Code plus text AI.
    - Rationale: Structured interface, side-by-side comparison with U of A GenAI, and alignment with Python-based workflow development.
    - Data sensitivity: High for resume data.
    - Institutional approval: Pilot/authorized setting only until broader approval is established.

Section 4 — Ethical Safeguards

- **Risk:** Exposure of sensitive applicant or employee data.

- Safeguard: Controlled pilot conditions, use of governed environments, and final records kept in HR systems rather than AI tools.
  - Responsible party: Compensation Manager, participating analysts, HR/IT reviewers as needed.
- **Risk:** Algorithmic bias in MQ or relevant experience analysis.
  - Safeguard: Human review of every output, fairness checks on sample cases, and documentation of overrides.
  - Responsible party: Participating analysts and manager.
- **Risk:** Over-reliance on AI in high-stakes decisions.
  - Safeguard: AI output is draft analysis only, never a final decision.
  - Responsible party: Compensation analysts and manager.
- **Risk:** Scope creep beyond MQ/relevant experience support.
  - Safeguard: Limit the workflow unless leadership authorizes added use cases.
  - Responsible party: Compensation Manager.

Section 5 — Equity Considerations

- **Population 1:** Applicants with non-traditional career paths, caregiving histories, disability-related accommodation history, international experience, or atypical job titles.
  - Adjustment: Explicitly look for equivalent experience and require manual review/documentation of edge cases.
- **Population 2:** Analysts with lower confidence using AI or a Streamlit interface.
  - Adjustment: Offer both workflow versions, written guidance, and hands-on support during the pilot.

Section 6 — Success Metrics

- **Metric 1:** Consistency of MQ/relevant experience conclusions.
  - How measured: Compare analyst-only, U of A GenAI-assisted, and Python/Streamlit-assisted reviews.
  - Timeline: During pilot, reviewed by Day 30–60.
  - Minimum acceptable threshold: At least a 10 percentage-point improvement in agreement or clear reduction in unexplained variation.
- **Metric 2:** Documentation quality and verification completeness.
  - How measured: Review pilot files for complete inputs, AI use notes, and human verification.
  - Timeline: Throughout pilot.

- Minimum acceptable threshold: At least 90% of pilot cases fully documented.

- **Metric 3:** Usefulness and workflow trust.
  - How measured: Analyst ratings and feedback.
  - Timeline: End of pilot.
  - Minimum acceptable threshold: Average usefulness rating of at least 4 out of 5.
- **Metric 4:** Vigilance on tricky cases.
  - How measured: Seed known-tricky cases and track whether reviewers flag them.
  - Timeline: Day 30 and Day 60.
  - Minimum acceptable threshold: At least 80% of planted/tricky cases flagged with substantive comment.

Section 7 — Rollout Plan — First Three Steps

- **Step 1:** Finalize prompt and Streamlit alignment so both versions use the same logic and output structure.
  - Who: Gina, with one senior Compensation analyst reviewing.
  - Resource needed: U of A GenAI access, Claude API access, sample resumes, JDXpert job descriptions.
  - By what date: Within 2 weeks of course completion.
- **Step 2:** Run a small AI-assisted pilot on resumes and job descriptions, preferably starting with low-risk, de-identified, or synthetic cases if needed.
  - Who: Gina and participating analysts.
  - Resource needed: Pilot cases, tracking sheet, guidance notes.
  - By what date: Weeks 3–5.
- **Step 3:** Prepare a short summary for the Compensation Manager with findings and recommendation.
  - Who: Gina.
  - Resource needed: Pilot data, analyst feedback, guideline draft.
  - By what date: By week 6 / within 30–45 days.

50-word elevator pitch

This pilot tests two AI-assisted workflows: (1) U of A GenAI and (2) a Python/Streamlit interface to support MQ and relevant experience review on complex resumes. The goal is better consistency, documentation, and fairness, while keeping all final judgments with Compensation analysts and using strong human verification at every step.