



RE: Module 6 Task 3 - AI Adoption Initiative Design Canvas

From Bishop, Jeff - (jeffbis) <jeffbis@arizona.edu>

Date Thu 2026-07-16 10:45 AM

To Jones, Gina - (ginajones) <ginajones@arizona.edu>

Gina,

Thank you so much for sending both of these over - and especially for the honesty about which one you actually like. That kind of candor makes for a much better trade, and I loved getting to spend real time with your work. This was genuinely fun to read.

On the MQ / Relevant Experience Review Pilot (the one you like - and you should!)

I want to start by just saying: this is a strong, thoughtful initiative, and it's strong in exactly the way our Module 6 rubric is looking for, not just "well-organized on the surface." A couple of things made me smile as I read through it:

Your equity section is honestly the best paragraph in either document, and maybe one of the best I've read in this cohort. Naming applicants with career breaks, caregiving histories, international experience, and unusual titles as the group most likely to be flattened by an automated read - and then giving the workflow an actual instruction to look for equivalent experience instead of just promising to "be fair" - is exactly the move Module 4 kept pushing all of us toward. Bias usually shows up as a missing consideration, not a wrong answer, and you found the missing consideration before anyone had to point it out to you. That's the kind of instinct that's hard to teach.

I also really appreciated how honest your rollout plan is about what you already have. Two weeks, six weeks, your own review, tools you're already cleared to use - nothing in Steps 1-3 is waiting on some hypothetical future approval. That's rarer than it sounds. A lot of Module 6 canvases (mine included, in an earlier draft!) quietly describe a future state as if it already exists. Yours doesn't do that, and it makes the whole plan feel genuinely doable, not just aspirational.

Here's the one place I'd gently push back, offered with real respect for how much you've already thought through: your safeguard against bias right now is "a human reviews every output." That's necessary, but on its own it's not quite enough, because the most common failure mode of human-in-the-loop review isn't a careless reviewer - it's a good reviewer who, after forty cases where the AI draft matched their own judgment, quietly stops re-deriving the answer and starts just confirming it. Your Section 6 metric (85% of cases showing no material change after human review) would look identical whether your reviewers are catching real errors or gently rubber-stamping. One idea: a small number of planted or known-tricky cases seeded into the pilot, to check whether the human catch rate holds up over time - the same lesson Module 2 taught us about not trusting a tool's fluency as a stand-in for its accuracy. Your reviewers deserve the same generous skepticism, not because they aren't excellent at their jobs, but because vigilance decays quietly and predictably in exactly this kind of repetitive review work. Building in a check for that makes the workflow stronger, not less trusting of your team.

One more small thing, offered warmly as an accessibility person rather than a strict reviewer: if this pilot ever touches disability-related accommodation history - and career-break and non-linear-path applicants sometimes do have that in their file - it might be worth folding explicitly into the equity language you already wrote, rather than assuming "non-traditional career path" fully covers it. Not a new section, just a sentence. You've already written the right kind of sentence twice here; a third would be easy, and it would make an already-strong section even more complete.

On the Tracker Clean-Up / Dashboard Helper (the "backup")

I want to be really clear about this one: please don't be embarrassed by it, and don't assume "backup" means "weaker" - because it isn't. What you actually did here is apply Module 5's governance lesson before anyone graded you on it. Our cohort's comparative-policy work found that real institutions almost never build in a forced reason to slow down and wait for controls - they publish a policy and move on. Your canvas does the opposite, on purpose: it refuses to touch real compensation data until masked or synthetic data and leadership authorization exist, and it treats "the honest answer might be not yet" as a perfectly acceptable outcome of the pilot rather than a failure of it. That's a genuinely mature governance instinct, not a smaller idea, and it deserves to be seen that way. If your team is more comfortable starting here, you're not settling for the easier project - you're choosing the one with the shortest, safest distance between what the canvas promises and what's actually okay to do this month. That's worth being proud of.

Which one to bring to your team

I can't make that call for you, but here's a reframe that might help: it doesn't have to be either/or. The caution built into your second canvas could be a lovely first phase of your first one - prove the workflow is trustworthy on low-stakes, low-sensitivity ground before it ever touches an MQ decision that shapes a real person's employment outcome. If your team is nervous about canvas one, that sequencing might be exactly the answer to their nervousness, rather than a reason to shelve the idea you actually believe in.

Truly, thank you for sharing this with me - it's clear how much care went into both of these, and I really enjoyed the chance to think alongside your work. Sending mine over separately now that you've so generously gone first. Looking forward to comparing notes!

Stay tuned!

Warmly,
Jeff

From: Jones, Gina - (ginajones) <ginajones@arizona.edu>
Sent: Thursday, July 16, 2026 9:50 AM
To: Bishop, Jeff - (jeffbis) <jeffbis@arizona.edu>
Subject: Module 6 Task 3 - AI Adoption Initiative Design Canvas

Jeff,

I'm attaching two because I don't think my team's going to like the first initiative that I did (even though I like it). You don't have to review them both.

Thank you and please send me yours when you're ready.

Gina Jones

Compensation Analyst

ginajones@arizona.edu

University of Arizona

Human Resources

compensation@arizona.edu



I really like to

ACCELERATE